

Goldman Sachs - Hackerrank - Car Popularity

Bala Kesavan

February 10, 2018

Context

To predict the popularity of a car given potential explanatory variables.

Exploratory data analysis

There are certain inferences we can draw about lack of popularity. Cars that are high on 'buying_price' or 'maintenance_cost' don't make the cut. Similarly, cars low on 'number_of_seats' or 'safety_rating' are not popular. 'number_of_doors' don't seem to make a difference while 'luggage_boot_size' has a weak effect on popularity.

The response variable, 'popularity', is ordinal. This has been taken care of in the pre-processing and in algorithm selection.

Next, several explanatory variables, like 'buying_price', are also ordinal. The choice was made to treat these as continuous variables after studying the well laid out arguments in this paper: Paper 248-2009, "Learning When to Be Discrete: Continuous vs. Categorical Predictors," David J. Pasta, ICON Clinical Research, San Francisco, CA.

Popularity levels 3 and 4 are under represented in the dataset. Since there are only 44 and 40 records each, a variety of subsampling techniques were tried, including SMOTE. Finally, 'upsample' was chosen since it performed best. This is understandable since SMOTE reduced the total count of data and that made it harder for the algorithm to learn.

```
library(caret)
library(ggplot2)

train_GS <- read.csv('train.csv')
summary(train_GS)
```

##	buying_price	maintenance_cost	number_of_doors	number_of_seats
##	Min. :1.000	Min. :1.000	Min. :2.000	Min. :2.000
##	1st Qu.:2.000	1st Qu.:2.000	1st Qu.:2.000	1st Qu.:2.000
##	Median :3.000	Median :3.000	Median :3.000	Median :4.000
##	Mean :2.533	Mean :2.528	Mean :3.494	Mean :3.633
##	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.250	3rd Qu.:5.000
##	Max. :4.000	Max. :4.000	Max. :5.000	Max. :5.000
##	luggage_boot_size	safety_rating	popularity	
##	Min. :1.000	Min. :1.000	Min. :1.000	
##	1st Qu.:1.000	1st Qu.:1.000	1st Qu.:1.000	
##	Median :2.000	Median :2.000	Median :1.000	

```
## Mean      :1.987      Mean      :1.978      Mean      :1.348
## 3rd Qu.:3.000      3rd Qu.:3.000      3rd Qu.:2.000
## Max.      :3.000      Max.      :3.000      Max.      :4.000
```

```
str(train_GS)
```

```
## 'data.frame': 1628 obs. of 7 variables:
## $ buying_price : int 3 3 1 4 3 2 1 3 1 2 ...
## $ maintenance_cost: int 2 2 4 4 3 1 3 1 1 1 ...
## $ number_of_doors : int 4 2 2 2 3 2 5 2 3 4 ...
## $ number_of_seats : int 2 5 5 2 4 2 2 4 5 2 ...
## $ luggage_boot_size: int 2 2 1 1 3 1 2 3 2 2 ...
## $ safety_rating : int 2 1 3 2 3 1 2 2 1 2 ...
## $ popularity : int 1 1 1 1 2 1 1 2 1 1 ...
```

```
train_GS_proc <- train_GS
```

```
train_GS_proc$popularity <- as.ordered(train_GS$popularity)
```

```
str(train_GS_proc)
```

```
## 'data.frame': 1628 obs. of 7 variables:
## $ buying_price : int 3 3 1 4 3 2 1 3 1 2 ...
## $ maintenance_cost: int 2 2 4 4 3 1 3 1 1 1 ...
## $ number_of_doors : int 4 2 2 2 3 2 5 2 3 4 ...
## $ number_of_seats : int 2 5 5 2 4 2 2 4 5 2 ...
## $ luggage_boot_size: int 2 2 1 1 3 1 2 3 2 2 ...
## $ safety_rating : int 2 1 3 2 3 1 2 2 1 2 ...
## $ popularity : Ord.factor w/ 4 levels "1"<"2"<"3"<"4": 1 1 1 1 2 1
1 2 1 1 ...
```

```
lapply(train_GS[, names(train_GS)], table)
```

```
## $buying_price
##
## 1 2 3 4
## 384 408 421 415
##
## $maintenance_cost
##
## 1 2 3 4
## 392 404 412 420
##
## $number_of_doors
##
## 2 3 4 5
## 411 409 401 407
##
## $number_of_seats
##
## 2 4 5
## 565 530 533
##
```

```

## $luggage_boot_size
##
##   1   2   3
## 553 543 532
##
## $safety_rating
##
##   1   2   3
## 565 534 529
##
## $popularity
##
##   1     2     3     4
## 1185  359   44   40

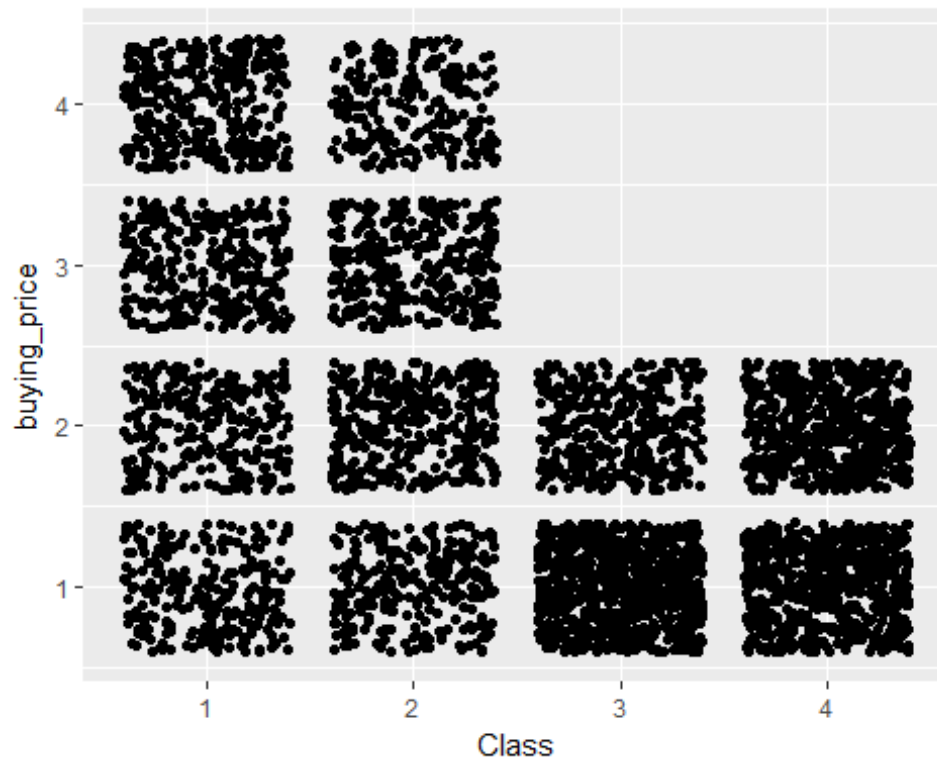
#(ftable(xtabs(popularity~buying_price+maintenance_cost , data = train_GS)))

#fixing the severe under representation of popularity levels 3 & 4
train_GS_proc_upsample <- upSample(x=train_GS_proc[, -7], y=
train_GS_proc$popularity)
table(train_GS_proc_upsample$Class)


##
##   1     2     3     4
## 1185 1185 1185 1185

#visual checks of informativeness/ ability to discriminate the response
variable
ggplot(train_GS_proc_upsample, aes(Class, buying_price)) + geom_jitter()


```




```
ggplot(train_GS_proc_upsample, aes(Class, maintenance_cost)) + geom_jitter()
```



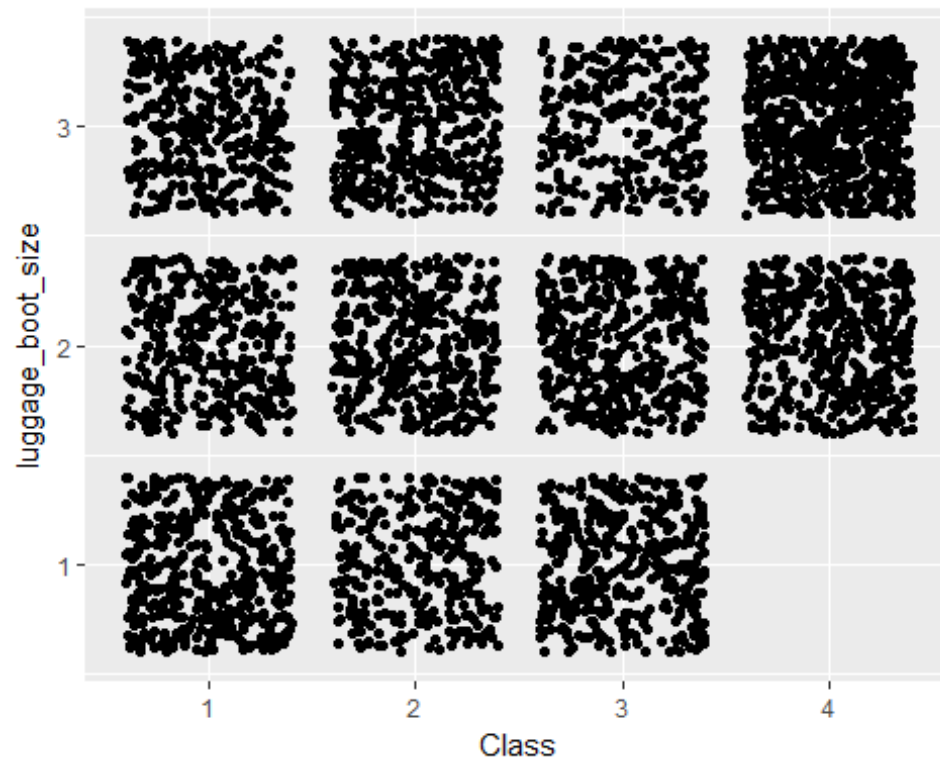
```
ggplot(train_GS_proc_upsample, aes(Class, number_of_doors)) + geom_jitter()
```



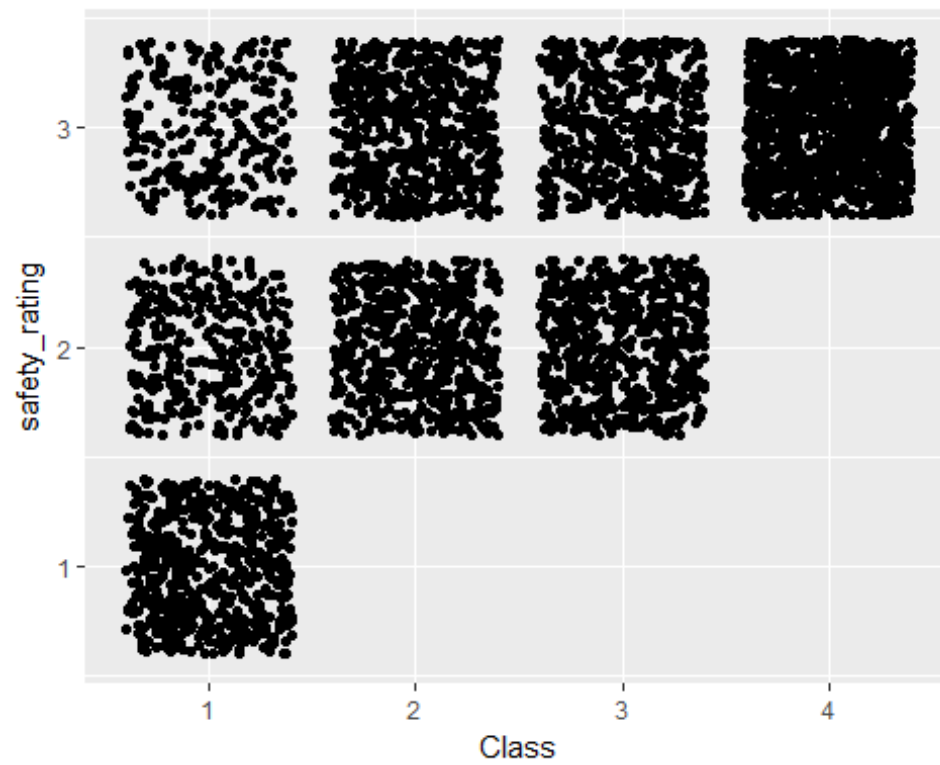
```
ggplot(train_GS_proc_upsample, aes(Class, number_of_seats)) + geom_jitter()
```



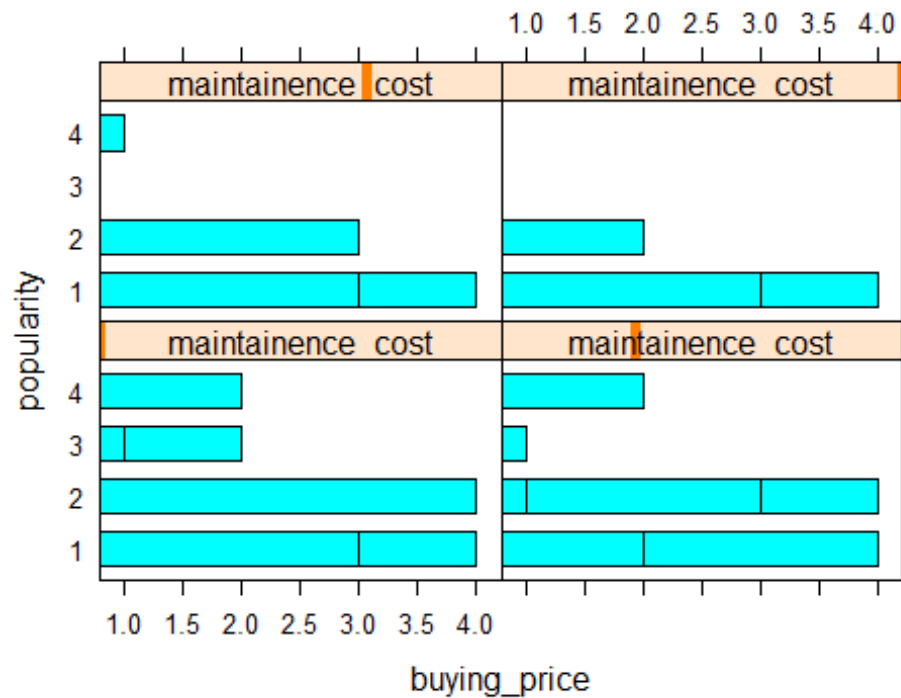
```
ggplot(train_GS_proc_upsample, aes(Class, luggage_boot_size)) + geom_jitter()
```



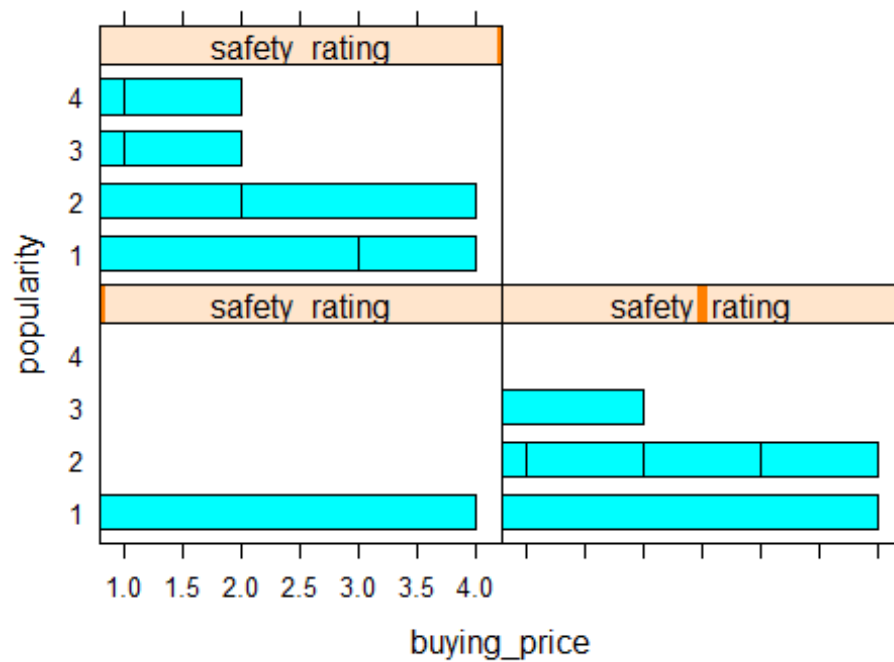
```
ggplot(train_GS_proc_upsample, aes(Class, safety_rating)) + geom_jitter()
```



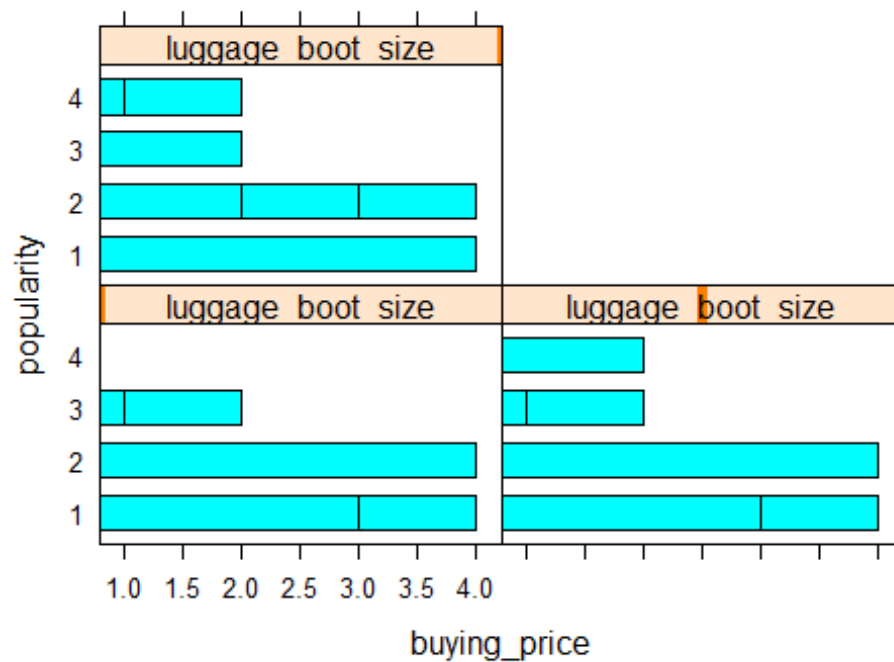
```
barchart(Class ~ buying_price | maintenance_cost, data =
train_GS_proc_upsample,
        xlab="buying_price", ylab="popularity")
```



```
barchart(Class ~ buying_price | safety_rating, data = train_GS_proc_upsample,
        xlab="buying_price", ylab="popularity")
```



```
barchart(Class ~ buying_price | luggage_boot_size, data =
train_GS_proc_upsample,
xlab="buying_price", ylab="popularity")
```




```

#statistical tests of independence
chisq.test(table(as.factor(train_GS$number_of_doors), train_GS$popularity))

##
## Pearson's Chi-squared test
##
## data:  table(as.factor(train_GS$number_of_doors), train_GS$popularity)
## X-squared = 8.0269, df = 9, p-value = 0.5314

chisq.test(table(as.factor(train_GS$number_of_seats), train_GS$popularity))

##
## Pearson's Chi-squared test
##
## data:  table(as.factor(train_GS$number_of_seats), train_GS$popularity)
## X-squared = 323.68, df = 6, p-value < 2.2e-16

```

Modeling

Simple models are usually best and the 'ordinal' model with the logistic link has been used. The model coefficients are along expected lines. e.g. 'popularity' and 'buying_price' have an inverse relationship while there is a positive reeationship to 'safety_rating'. Several options were tested, like polynomial and interaction terms. The best performing option, with one interaction term, was determined using the ANOVA test as shown below.

```

library(ordinal)
model_clm_GS_upsample <-
clm(Class~buying_price+maintainence_cost+luggage_boot_size+safety_rating,
data = train_GS_proc_upsample)
summary(model_clm_GS_upsample)

## formula:
## Class ~ buying_price + maintainence_cost + luggage_boot_size +
safety_rating
## data:    train_GS_proc_upsample
##
## link threshold nobis logLik   AIC      niter max.grad cond.H
## logit flexible  4740 -3611.70 7237.40 6(0)  1.37e-08 2.7e+03
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## buying_price    -1.77383    0.04510  -39.33  <2e-16 ***
## maintainence_cost -1.32913    0.04047  -32.84  <2e-16 ***
## luggage_boot_size  1.64197    0.04816   34.10  <2e-16 ***
## safety_rating      3.42819    0.07700   44.52  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Threshold coefficients:

```

```
##      Estimate Std. Error z value
## 1|2    2.9580     0.2034  14.54
## 2|3    6.0688     0.2154  28.17
## 3|4    8.6892     0.2370  36.66

model_clm_GS_upsample1 <- update(model_clm_GS_upsample, ~. +
  buying_price*safety_rating)

anova(model_clm_GS_upsample, model_clm_GS_upsample1)

## Likelihood ratio tests of cumulative link models:
##
##              formula:
## model_clm_GS_upsample  Class ~ buying_price + maintenance_cost +
  luggage_boot_size + safety_rating
## model_clm_GS_upsample1 Class ~ buying_price + maintenance_cost +
  luggage_boot_size + safety_rating + buying_price:safety_rating
##              link: threshold:
## model_clm_GS_upsample  logit flexible
## model_clm_GS_upsample1 logit flexible
##
##              no.par      AIC  logLik LR.stat df Pr(>Chisq)
## model_clm_GS_upsample      7 7237.4 -3611.7
## model_clm_GS_upsample1     8 7088.2 -3536.1  151.16  1 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

test_GS <- read.csv('test.csv')
names(test_GS) <-
c("buying_price", "maintenance_cost", "number_of_doors", "number_of_seats", "luggage_boot_size", "safety_rating")

prediction <- predict(model_clm_GS_upsample1, newdata = test_GS, type =
  'class')

write.csv(prediction, file = 'prediction.csv', row.names = FALSE)
```

Conclusion

The best score obtained during submissions by Hackerrank's checker was:

F1 score = [1: 0.181818, 2: 0.260870, 3: 0.444444, 4: 0.333333]. In a way, the better performance for levels 3 and 4 is to be expected. It is easier to model the fact that cars cannot reach levels 3 or 4 of popularity without performing very well on 'buying_price', 'maintenance_cost' and 'safety_rating'.

But the reverse isn't true. Just because a car performs very well on 'buying_price', 'maintenance_cost' and 'safety_rating' doesn't lead to higher popularity.

Finding this out could be the next step.

Appendix showing confusion matrices and demonstrating the use of other algorithms

More algorithms are available in caret.

```
#data partitioning
set.seed(123)
inTraining <- createDataPartition(train_GS_proc_upsample$Class, p = 0.8, list
= FALSE)
training <- train_GS_proc_upsample[inTraining, ]
testing <- train_GS_proc_upsample[-inTraining, ]

#modeling with CLM
library(ordinal)
#model_clm_GS_upsample <-
clm(Class~buying_price+maintenance_cost+luggage_boot_size+safety_rating,
data = training)
model_clm_GS_upsample <- readRDS('model_clm_GS_upsample.rds')
summary(model_clm_GS_upsample)

## formula:
## Class ~ buying_price + maintenance_cost + luggage_boot_size +
safety_rating
## data:   training
##
## link threshold nobs logLik   AIC      niter max.grad cond.H
## logit flexible  3792 -2918.81 5851.62 6(0)  5.87e-09 2.7e+03
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## buying_price      -1.71273    0.04997  -34.27  <2e-16 ***
## maintenance_cost  -1.32530    0.04477  -29.60  <2e-16 ***
## luggage_boot_size   1.65398    0.05443   30.39  <2e-16 ***
## safety_rating       3.33787    0.08454   39.48  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Threshold coefficients:
##      Estimate Std. Error z value
## 1|2    2.9690    0.2302  12.90
## 2|3    6.0286    0.2424  24.87
## 3|4    8.6054    0.2658  32.38

#saveRDS(model_clm_GS_upsample, file='model_clm_GS_upsample.rds')
#model_clm_GS_upsample1 <- update(model_clm_GS_upsample, ~. +
buying_price*safety_rating)
model_clm_GS_upsample1 <- readRDS('model_clm_GS_upsample1.rds')
summary(model_clm_GS_upsample1)
```

```

## formula:
## Class ~ buying_price + maintenance_cost + luggage_boot_size +
safety_rating + buying_price:safety_rating
## data:   training
##
## link threshold nobs logLik   AIC      niter max.grad cond.H
## logit flexible 3792 -2863.11 5742.23 6(0) 5.20e-09 2.0e+04
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## buying_price      0.05637    0.16453   0.343   0.732
## maintenance_cost  -1.35113    0.04520 -29.891 <2e-16 ***
## luggage_boot_size   1.69395    0.05525  30.659 <2e-16 ***
## safety_rating       4.71824    0.16072  29.357 <2e-16 ***
## buying_price:safety_rating -0.70803    0.06505 -10.884 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Threshold coefficients:
##      Estimate Std. Error z value
## 1|2    6.3666    0.4012  15.87
## 2|3    9.4468    0.4126  22.90
## 3|4   12.1727    0.4425  27.51

anova(model_clm_GS_upsample, model_clm_GS_upsample1)

## Likelihood ratio tests of cumulative link models:
##
##              formula:
## model_clm_GS_upsample Class ~ buying_price + maintenance_cost +
luggage_boot_size + safety_rating
## model_clm_GS_upsample1 Class ~ buying_price + maintenance_cost +
luggage_boot_size + safety_rating + buying_price:safety_rating
##              link: threshold:
## model_clm_GS_upsample  logit flexible
## model_clm_GS_upsample1 logit flexible
##
##              no.par    AIC  logLik LR.stat df Pr(>Chisq)
## model_clm_GS_upsample      7 5851.6 -2918.8
## model_clm_GS_upsample1     8 5742.2 -2863.1  111.39  1 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

#saveRDS(model_clm_GS_upsample1, file='model_clm_GS_upsample1.rds')
prediction <- predict(model_clm_GS_upsample1, newdata = testing, type =
'class')

confusionMatrix(as.numeric(prediction$fit), as.numeric(testing$Class))

## Confusion Matrix and Statistics
##

```

```

##           Reference
## Prediction   1   2   3   4
##           1 172  55   0   0
##           2  40 136  38   0
##           3  16  46 146  34
##           4   9   0  53 203
##
## Overall Statistics
##
##           Accuracy : 0.693
##           95% CI : (0.6626, 0.7223)
##           No Information Rate : 0.25
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.5907
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity      0.7257   0.5738   0.6160   0.8565
## Specificity      0.9226   0.8903   0.8650   0.9128
## Pos Pred Value   0.7577   0.6355   0.6033   0.7660
## Neg Pred Value   0.9098   0.8624   0.8711   0.9502
## Prevalence       0.2500   0.2500   0.2500   0.2500
## Detection Rate   0.1814   0.1435   0.1540   0.2141
## Detection Prevalence 0.2395   0.2257   0.2553   0.2795
## Balanced Accuracy 0.8242   0.7321   0.7405   0.8847

```

#modeling with 'CART or Ordinal Responses'

```

Control <- trainControl(method = 'cv', 10)
set.seed(21)
#model_GS_rpartScore <- train(Class~., data = train_GS_proc_upsample,
method="rpartScore", trControl=Control)
(model_GS_rpartScore <- readRDS('model_GS_rpartScore.rds'))

```

```

## CART or Ordinal Responses
##
## 4740 samples
## 6 predictor
## 4 classes: '1', '2', '3', '4'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 4267, 4266, 4267, 4266, 4267, 4267, ...
## Resampling results across tuning parameters:
##
##  cp          split  prune  Accuracy  Kappa
##  0.1136428  abs    mr     0.5647882 0.4196945
##  0.1136428  abs    mc     0.3905061 0.1873336

```

```
## 0.1136428 quad mr 0.5261631 0.3682044
## 0.1136428 quad mc 0.3905061 0.1873336
## 0.1477731 abs mr 0.5261631 0.3682044
## 0.1477731 abs mc 0.3905061 0.1873336
## 0.1477731 quad mr 0.5261631 0.3682044
## 0.1477731 quad mc 0.3905061 0.1873336
## 0.1589311 abs mr 0.5261631 0.3682044
## 0.1589311 abs mc 0.3905061 0.1873336
## 0.1589311 quad mr 0.5261631 0.3682044
## 0.1589311 quad mc 0.3905061 0.1873336
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were cp = 0.1136428, split = abs
## and prune = mr.
```

```
#saveRDS(model_GS_rpartScore, file = 'model_GS_rpartScore.rds')
prediction <- predict(model_GS_rpartScore, newdata = testing)
```

```
confusionMatrix((prediction), (testing$Class))
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction  1    2    3    4
##           1 139 119    0    0
##           2   0   0   0   0
##           3  79  68 128   0
##           4  19  50 109 237
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.5316
##           95% CI : (0.4993, 0.5638)
##           No Information Rate : 0.25
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##           Kappa : 0.3755
##           McNemar's Test P-Value : < 2.2e-16
```

```
##
## Statistics by Class:
```

```
##
##           Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity      0.5865      0.00      0.5401      1.0000
## Specificity      0.8326      1.00      0.7932      0.7496
## Pos Pred Value    0.5388      NaN      0.4655      0.5711
## Neg Pred Value    0.8580      0.75      0.8380      1.0000
## Prevalence       0.2500      0.25      0.2500      0.2500
## Detection Rate    0.1466      0.00      0.1350      0.2500
## Detection Prevalence 0.2722      0.00      0.2901      0.4378
## Balanced Accuracy 0.7096      0.50      0.6667      0.8748
```

```

#modeling with 'Ordered Logistic or Probit Regression'
Control <- trainControl(method = 'cv', 10)
set.seed(21)
#model_GS_polr <- train(Class~., data = train_GS_proc_upsample,
method="polr", trControl=Control)
(model_GS_polr <- readRDS('model_GS_polr.rds'))

## Ordered Logistic or Probit Regression
##
## 4740 samples
##    6 predictor
##    4 classes: '1', '2', '3', '4'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 4265, 4266, 4267, 4266, 4266, 4264, ...
## Resampling results across tuning parameters:
##
##  method      Accuracy      Kappa
##  cauchit    0.7293173    0.6390929
##  cloglog           NaN           NaN
##  logistic   0.7248882    0.6331878
##  loglog     0.6862684    0.5816932
##  probit     0.7139114    0.6185500
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was method = cauchit.

#saveRDS(model_GS_polr, file = 'model_GS_polr.rds')

prediction <- predict(model_GS_polr, newdata = testing)

confusionMatrix((prediction), (testing$Class))

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    1    2    3    4
##           1 182   23    0    0
##           2  47  164   27    0
##           3   6   42  170   47
##           4   2    8   40  190
##
## Overall Statistics
##
##              Accuracy : 0.7447
##              95% CI : (0.7157, 0.7722)
##    No Information Rate : 0.25
##    P-Value [Acc > NIR] : < 2.2e-16
##

```

```

##              Kappa : 0.6596
## McNemar's Test P-Value : 9.184e-05
##
## Statistics by Class:
##
##              Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity      0.7679   0.6920   0.7173   0.8017
## Specificity      0.9677   0.8959   0.8664   0.9297
## Pos Pred Value   0.8878   0.6891   0.6415   0.7917
## Neg Pred Value   0.9260   0.8972   0.9019   0.9336
## Prevalence       0.2500   0.2500   0.2500   0.2500
## Detection Rate   0.1920   0.1730   0.1793   0.2004
## Detection Prevalence 0.2162   0.2511   0.2795   0.2532
## Balanced Accuracy 0.8678   0.7940   0.7918   0.8657

#modeling with 'Adjacent Categories Probability Model for Ordinal Data'
Control <- trainControl(method = 'cv', 10)
set.seed(21)
#model_GS_vglmAdjCat <- train(Class~., data = train_GS_proc_upsample,
method="vglmAdjCat", trControl=Control)
(model_GS_vglmAdjCat <- readRDS('model_GS_vglmAdjCat.rds'))

## Adjacent Categories Probability Model for Ordinal Data
##
## 4740 samples
## 6 predictor
## 4 classes: '1', '2', '3', '4'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 4265, 4266, 4267, 4266, 4266, 4264, ...
## Resampling results across tuning parameters:
##
##  parallel  Accuracy  Kappa
##  FALSE     0.8750812  0.8334416
##  TRUE      0.7250956  0.6334630
##
## Tuning parameter 'link' was held constant at a value of loge
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were parallel = FALSE and link = loge.

#saveRDS(model_GS_vglmAdjCat, file = 'model_GS_vglmAdjCat.rds')
prediction <- predict(model_GS_vglmAdjCat, newdata = testing)

confusionMatrix((prediction), (testing$Class))

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    1    2    3    4
##              1 200 218 157 176

```



```

##          2  25   8  80  25
##          3  12  11   0  36
##          4   0   0   0   0
##
## Overall Statistics
##
##              Accuracy : 0.2194
##              95% CI : (0.1934, 0.2471)
##      No Information Rate : 0.25
##      P-Value [Acc > NIR] : 0.9874
##
##              Kappa : -0.0408
##  McNemar's Test P-Value : <2e-16
##
## Statistics by Class:
##
##              Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity          0.8439 0.033755 0.00000    0.00
## Specificity          0.2250 0.817159 0.91702    1.00
## Pos Pred Value       0.2663 0.057971 0.00000    NaN
## Neg Pred Value       0.8122 0.717284 0.73341    0.75
## Prevalence           0.2500 0.250000 0.25000    0.25
## Detection Rate       0.2110 0.008439 0.00000    0.00
## Detection Prevalence 0.7922 0.145570 0.06224    0.00
## Balanced Accuracy     0.5345 0.425457 0.45851    0.50

```